

Untitled CAV Paper

Alexander Legg¹, Leonid Ryzhyk², and Nina Narodytska²

¹ NICTA* and UNSW

`alexander.legg@nicta.com.au`

² Samsung Research

Abstract. Reactive synthesis techniques based on image computations have been shown to work well in many cases but suffer from state explosion in others. A different approach applies SAT solvers in a counterexample guided framework but is limited to bounded synthesis. We present an extension of this technique to unbounded synthesis by employing interpolation. Experimental results show this to be a promising alternative that solves some previously intractable instances.

1 Introduction

Reactive systems are ubiquitous in real-world problems such as circuit design, industrial automation, or device drivers. Automatic synthesis can provide a *correct by construction* controller for a reactive system from a specification. However, the reactive synthesis problem is 2EXPTIME-complete so naive algorithms are infeasible on even simple systems.

Reactive synthesis is formalised as a game between the *controller* and its *environment*. In this work we focus on safety games, in which the controller must prevent the environment from forcing the game into an error state. Much of the complexity of reactive synthesis stems from tracking the set of states in which a player is winning.

There are several techniques that aim to mitigate this complexity by representing states symbolically. Historically the most successful technique has been to use *Binary Decision Diagrams* (BDDs). BDDs efficiently represent a relation on a set of game variables but in the worst case the representation may be exponential. This means that BDDs are not a one-size-fits-all solution for all reactive synthesis specifications.

Advances in SAT solving technology has prompted research into its applicability to synthesis as an alternative to BDDs. One approach is to find sets of states in CNF [?]. Another approach is to eschew states and focus on *runs* of the game. Previous work has applied this idea to realizability of bounded games [3] by forming abstract representations of the game. In this paper, we extend this idea to unbounded games by constructing approximate sets of winning states from abstract trees.

* NICTA is funded by the Australian Government as represented by the Department of Broadband, Communications and the Digital Economy and the Australian Research Council through the ICT Centre of Excellence program.

2 Reactive Synthesis

A *safety game*, $G = \langle X, L_u, L_c, \delta, I, E, \rangle$, consists of a set of boolean state variables, sets of uncontrollable and controllable label variables, a transition relationship $\delta : (X, L_u, L_c) \rightarrow X$, an initial state, and an error state. The *controller* and *environment* players choose controllable and uncontrollable labels respectively and the game proceeds according to δ .

An *run* of a game $(x_0, u_0, c_0), (x_1, u_1, c_1) \dots (x_n, u_n, c_n)$ is a chain of state and label pairs of length n s.t. $x_{k+1} \leftarrow \delta(x_k, u_k, c_k)$. A run is winning for the controller if $x_0 = I \wedge \forall i \in \{1..n\} (x_i \neq E)$. In a bounded game of rank n all runs are restricted to length n , whereas unbounded games consider runs of infinite length. Since we consider only deterministic games a run is equivalent to a list of assignments to L_c and L_u .

A *controller strategy* $\pi^c : (X, L_u) \rightarrow L_c$ is a mapping of states and uncontrollable inputs to controllable labels. A controller strategy is winning in a bounded game of rank n if all runs $(x_0, u_0, \pi^c(x_0, u_0)), (x_1, u_1, \pi^c(x_1, u_1)) \dots (x_n, u_n, \pi^c(x_n, u_n))$ are winning. Bounded *realizability* is the problem of determining the existence of such a strategy for a bounded game.

An *environment strategy* $\pi^e : X \rightarrow L_u$ is a mapping of states to uncontrollable labels. A bounded run is winning for the environment if $x_0 = I \wedge \exists i \in \{1..n\} (x_i = E)$ and an environment strategy is winning for a bounded game if there exists a run $(x_0, \pi^e(x_1), c_1), (x_1, \pi^e(x_1), c_1) \dots (x_n, \pi^e(x_n), c_n)$ that wins for the environment. Safety games are zero sum, therefore the existence of a controller strategy implies the nonexistence of an environment strategy and vice versa.

2.1 Abstract Game Trees

A set of runs can be represented symbolically by a tree of valuations to controllable and uncontrollable label variables where valuations are allowed to be *unfixed*. An unfixed value symbolically represents all possible valuations. Formally, we define the tree as a set of lists of label valuations and a node in the tree may be identified by the list made by following the path from the root.

The entire set of runs of a bounded game of rank n is symbolically represented by a tree of depth $2n$ and width 1 populated by unfixed edges. Reducing the set of runs in a game forms an abstract game, which can be represented symbolically by an *abstract game tree*.

A strategy is equivalent to the set of all runs with the player's labels obeying the strategy mapping π . Therefore, a strategy can also be represented by a tree. Relaxing the restriction of the strategy mapping allows for a *partial strategy* in which multiple labels are now possible for a single state.

A strategy or partial strategy can also be thought of as an abstract game. A partial strategy for the controller is a restriction only on the controllable labels in the game. So if the environment can not win in the abstract game equivalent to the controller's partial strategy, then all strategies allowable by that partial strategy must be winning.

An abstract game tree can be checked for the existence of a winning run by a SAT solver[3]. The tree must be encoded into CNF by making copies of the transition relation for each game step in tree. The formula must also check whether the error state has been reached in no branch/in all branches for the controller and environment respectively. The functions that produce these formulas are shown in Algorithm 2.

Algorithm 1 Tree formulas for Controller and Environment respectively

```

function TREEFORMULA(gt)
  if RANK(gt) == 0 then
    return  $\neg E(x^{gt})$ 
  else
    return  $\neg E(x^{gt}) \wedge \bigwedge_{n \in \text{SUCC}(gt)} (\delta(n) \wedge \text{LABEL}(n) \wedge \text{TREEFORMULA}(n))$ 
  end if
end function

function  $\overline{\text{treeFormula}}$ (gt)
  if RANK(gt) == 0 then
    return  $E(x^{gt})$ 
  else
    return  $E(x^{gt}) \vee \bigvee_{n \in \text{SUCC}(gt)} (\delta(n) \wedge \text{LABEL}(n) \wedge \overline{\text{treeFormula}}(n))$ 
  end if
end function

```

When we produce a formula using TREEFORMULA the SAT solver is playing on behalf on the controller. Any unfixed labels in the tree, controllable or uncontrollable, will be existentially quantified. This means that if there exists a way for the game to be solved when both players cooperate the SAT solver will find it. If no winning run exists in an abstract game even when the players are cooperating then there is definitely no winning run when the opponent is playing adversarily.

2.2 Counterexample Guided Bounded Synthesis

By checking for the existence of a winning run for the environment in an abstract game tree constructed from a partial controller strategy we are checking if the partial strategy is winning. If no spoiling run can be found then the partial strategy must always win. If a spoiling run is found then we then we have a counterexample that proves that the controller strategy is not winning. This forms the basis of a counterexample guided abstraction refinement framework that operates on candidate strategies.

The bounded synthesis algorithm described in [3] begins with an empty game as seen in Figure 1a. Initially we are playing on behalf of the environment because it chooses the first move in each step. The empty game is passed to the SAT solver, which searches for a candidate environment strategy. If a candidate is found (Figure 1b) then it is checked for a spoiling strategy by solving

for the controller. If no spoiling strategy exists, that means our candidate is a winning strategy and the algorithm terminates. Otherwise we find a counterexample (Figure 1c), which is used to refine the empty game tree to include the first move from the controller’s spoiling strategy (Figure 1d). The algorithm continues by finding a new candidate for the environment (Figure 1e), and a new counterexample to refine the game abstraction once again (Figure 1f).




Fig. 1: Abstract game trees

3 Unbounded Synthesis

Bounded synthesis can be used to prove the existence of a winning strategy for the environment by providing a witness. For the controller, the strongest claim that can be made is that the strategy is winning as long as the game does not extend longer than the bound.

In model checking the maximum bound is decided based on the states of the game [1]. The naïve approach is to use size of the state space as the bound ($2^{|X|}$). With a bound of this size all states may be explored. One optimisation is to use the diameter of the game, which is the smallest number d such that for any state x there is a path of length $\leq d$ to all other reachable states. However, for large games these bounds are intractable.

When performing model checking or synthesis with BDDs [2] the set of states that are winning for the environment is iteratively constructed by computing the states from which the environment can force the game into the previous winning set. Eventually this process reaches a fixed point and the total set of environment winning states is known.

A similar concept can be applied to the bounded synthesis algorithm to iteratively increase the bound of the game and terminate when a fixed point is

```

function SOLVEABSTRACT( $\langle s, r \rangle, gt$ )
   $cand \leftarrow \text{FINDCANDIDATE}(\langle s, r \rangle, gt)$ 
  if  $r = n - 1$  then return  $cand$ 
   $gt' \leftarrow gt$ 
  loop
    if  $cand = \emptyset$  then return  $\emptyset$ 
     $cex \leftarrow \text{VERIFY}(\langle s, r \rangle, gt, cand)$ 
    if  $cex = \text{NULL}$  then return  $cand$ 
     $gt' \leftarrow gt' \cup cex$ 
     $cand \leftarrow \text{SOLVEABSTRACT}(\langle s, r \rangle, gt')$ 
  end loop
end function

function FINDCANDIDATE( $\langle s, r \rangle, T$ )
   $\hat{T} \leftarrow \text{GT\_EXTEND}(T)$ 
   $fml \leftarrow \text{TREEFORMULA}(\hat{T})$ 
   $sol \leftarrow \text{SAT}(s(X_T) \wedge fml)$ 
end function

function VERIFY( $\langle r, s \rangle, gt, cand$ )
  for  $l \in \text{leaves}(gt)$  do
     $\langle r', s' \rangle \leftarrow \text{OUTCOME}(\langle r, s \rangle, gt, l)$ 
     $spoiling \leftarrow \text{SOLVEABSTRACT}(\langle r', s' \rangle, \emptyset)$ 
  end for
end function

```

reached. When a strategy is found to be winning on an abstract game tree, we record as winning the states from which the opponent could find no counterexample. To find these states we use interpolation of subformulas of the game tree.

3.1 Learning States with Interpolants

Given two formulas A and B such that $A \wedge B$ is unsatisfiable, it is possible to construct an interpolant \mathcal{I} such that $A \rightarrow \mathcal{I}$, $B \wedge \mathcal{I}$ is unsatisfiable, and \mathcal{I} refers only to the intersection of variables in A and B . An interpolant can be constructed efficiently from a resolution proof of the unsatisfiability of $A \wedge B$ [4].

Consider the snippet of a game tree in Figure 2. The tree is losing for the controller, the node labelled x_1 is at rank 1, and x_0^0 and x_0^1 are at rank 0. Since the tree is controller-losing we know that at least one run represented by the tree contains the error state. As a result, $\text{TREEFORMULA}(gt)$ is unsatisfiable. If we take the step x_1 to x_0^0 and cut it from the rest of the tree then $\text{TREEFORMULA}(step) \wedge \text{TREEFORMULA}(parent)$ must still be unsatisfiable.

We can construct an interpolant with $A = \text{TREEFORMULA}(parent)$ and $B = \text{TREEFORMULA}(step)$. The only variables shared between A and B are the state variables at x_1 . We know that $B \wedge \mathcal{I}$ is unsatisfiable, therefore all states in \mathcal{I} must lose to the uncontrollable label u_1 . We also know that $A \rightarrow \mathcal{I}$, thus \mathcal{I} contains all states reachable by the parent tree (on runs that avoid the error state.)




Fig. 2: A controller-losing game tree

Algorithm 2 Amended tree formulas for Controller and Environment respectively

```

function TREEFORMULA(gt)
  if RANK(gt) == 0 then
    return  $\neg L^c(x^{gt})$ 
  else
    return  $\neg L^c(x^{gt}) \wedge \bigwedge_{n \in \text{SUCC}(gt)} (\delta(n) \wedge \text{LABEL}(n) \wedge \text{TREEFORMULA}(n))$ 
  end if
end function

function  $\overline{\text{treeFormula}}$ (gt)
  if RANK(gt) == 0 then
    return  $E(x^{gt})$ 
  else
    return  $L^e[\text{rank}(gt)](x^{gt}) \wedge$ 
     $(E(x^{gt}) \vee \bigvee_{n \in \text{SUCC}(gt)} (\delta(n) \wedge \text{LABEL}(n) \wedge \overline{\text{treeFormula}}(n)))$ 
  end if
end function
  
```

Now we can consider the step x_1 to x_0^1 , and the parent - now without either $x_1 \rightarrow x_0^0$ or $x_1 \rightarrow x_0^1$. The formula $(\text{TREEFORMULA}(\text{parent}) \wedge \mathcal{I}) \wedge \text{TREEFORMULA}(\text{step})$ must be unsatisfiable. \mathcal{I} contains all states that lose to u_1 so any other state reachable at x_1 must lose to u_2 . Therefore we can compute another interpolant that contains states that lose to u_2 .

This is the foundation for a recursive algorithm that consumes an entire tree by removing a single step on each iteration. All learned states from which the controller must lose are recorded in a set L^c . This algorithm can also be performed on environment-losing trees with the caveat that any state learnt at a node of rank n is only known to lose at ranks less than or equal to n .

3.2 Proof of Termination

3.3 Algorithm

4 Evaluation

5 Related Work

6 Conclusion

References

1. Biere, A., Cimatti, A., Clarke, E., Zhu, Y.: Symbolic model checking without bdds. In: Cleaveland, W. (ed.) Tools and Algorithms for the Construction and Analysis of Systems, Lecture Notes in Computer Science, vol. 1579, pp. 193–207. Springer Berlin Heidelberg (1999)
2. Burch, J.R., Clarke, E.M., McMillan, K.L., Dill, D.L., Hwang, L.J.: Symbolic model checking: 10^{20} states and beyond. In: Symposium on Logic in Computer Science. pp. 428–439 (1990)
3. Narodytska, N., Legg, A., Bacchus, F., Ryzhyk, L., Walker, A.: Solving games without controllable predecessor. In: Computer Aided Verification. pp. 533–540 (2014)
4. Pudlák, P.: Lower bounds for resolution and cutting plane proofs and monotone computations. Journal of Symbolic Logic 62(3), 981–998 (1997)